

Introduction to Structured Query Language (SQL)



SQL

Yevhen Berkunskyi, NUoS
eugeny.berkunsky@gmail.com
<http://www.berkut.mk.ua>

Objectives

- Explore basic commands and functions of SQL
- How to use SQL for data administration (to create tables, indexes, and views)
- How to use SQL for data manipulation (to add, modify, delete, and retrieve data)
- How to use SQL to query a database to extract useful information



Introduction to SQL

- SQL functions fit into two broad categories:
 - Data definition language
 - SQL includes commands to:
 - Create database objects, such as tables, indexes, and views
 - Define access rights to those database objects
 - Data manipulation language
 - Includes commands to insert, update, delete, and retrieve data within database tables



Introduction to SQL (continued)

- SQL is relatively easy to learn
- Basic command set has vocabulary of less than 100 words
- Nonprocedural language
- American National Standards Institute (ANSI) prescribes a standard SQL
- Several SQL dialects exist



Introduction to SQL (continued)

SQL Data Definition commands

COMMAND OR OPTION	DESCRIPTION
CREATE SCHEMA AUTHORIZATION	Creates a database schema
CREATE TABLE	Creates a new table in the user's database schema
NOT NULL	Ensures that a column will not have null values
UNIQUE	Ensures that a column will not have duplicate values
PRIMARY KEY	Defines a primary key for a table
FOREIGN KEY	Defines a foreign key for a table
DEFAULT	Defines a default value for a column (when no value is given)
CHECK	Constraint used to validate data in an attribute
CREATE INDEX	Creates an index for a table
CREATE VIEW	Creates a dynamic subset of rows/columns from one or more tables
ALTER TABLE	Modifies a table's definition (adds, modifies, or deletes attributes or constraints)
CREATE TABLE AS	Creates a new table based on a query in the user's database schema
DROP TABLE	Permanently deletes a table (and thus its data)
DROP INDEX	Permanently deletes an index
DROP VIEW	Permanently deletes a view

Introduction to SQL (continued)

SQL Data Manipulation commands

COMMAND OR OPTION	DESCRIPTION
INSERT	Inserts row(s) into a table
SELECT	Selects attributes from rows in one or more tables or views
WHERE	Restricts the selection of rows based on a conditional expression
GROUP BY	Groups the selected rows based on one or more attributes
HAVING	Restricts the selection of grouped rows based on a condition
ORDER BY	Orders the selected rows based on one or more attributes
UPDATE	Modifies an attribute's values in one or more table's rows
DELETE	Deletes one or more rows from a table
COMMIT	Permanently saves data changes
ROLLBACK	Restores data to their original values

Introduction to SQL (continued)

SQL Data Manipulation commands (continued)

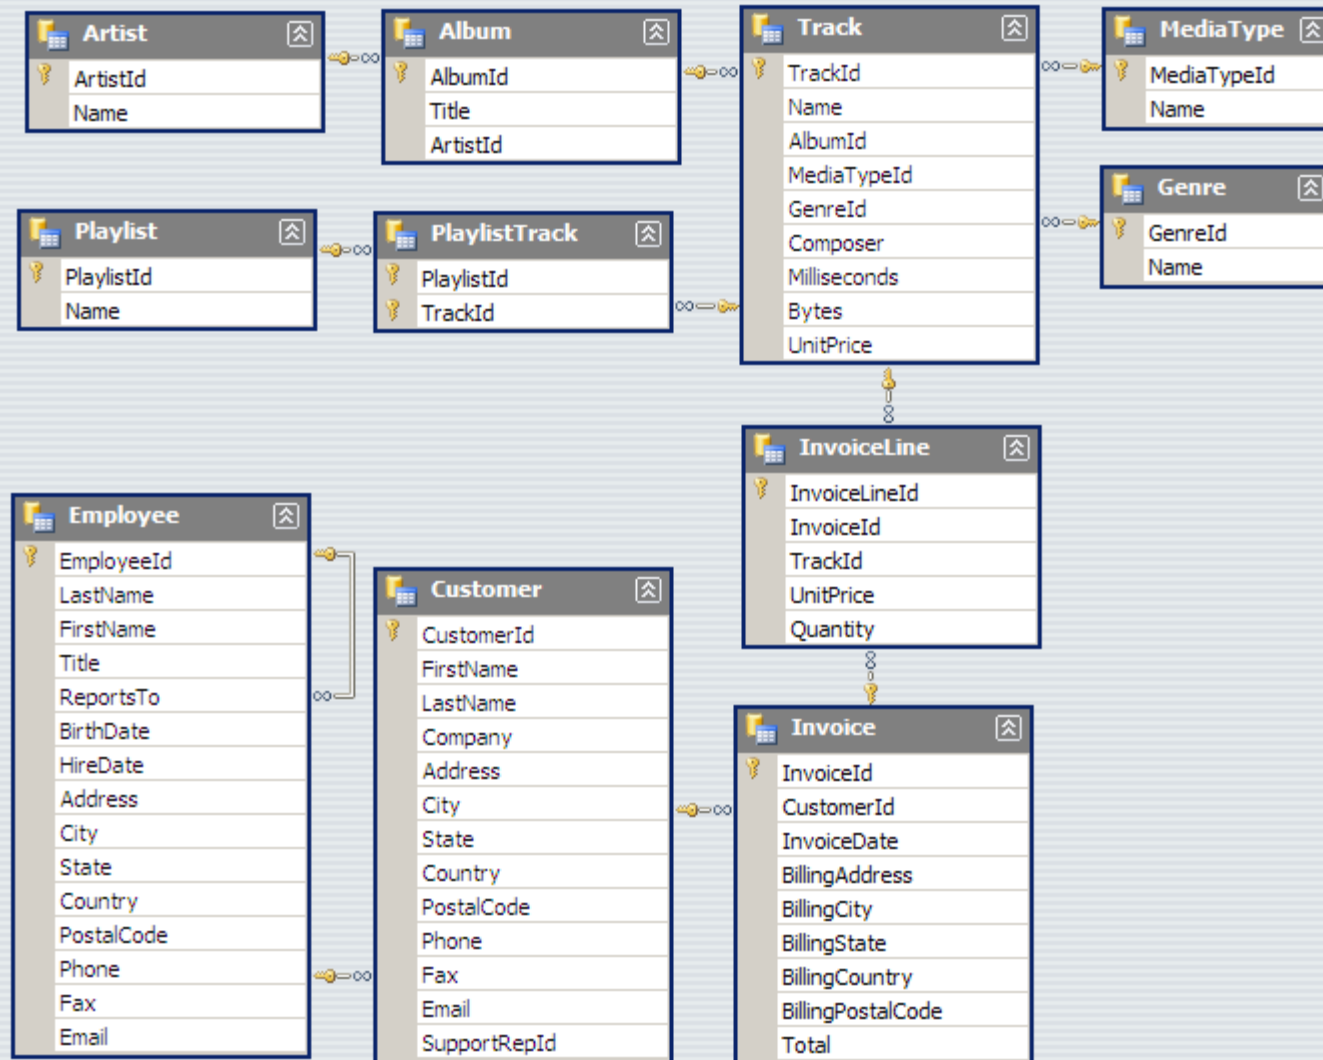
COMMAND OR OPTION	DESCRIPTION
COMPARISON OPERATORS	
=, <, >, <=, >=, <>	Used in conditional expressions
LOGICAL OPERATORS	
AND/OR/NOT	Used in conditional expressions
SPECIAL OPERATORS	Used in conditional expressions
BETWEEN	Checks whether an attribute value is within a range
IS NULL	Checks whether an attribute value is null
LIKE	Checks whether an attribute value matches a given string pattern
IN	Checks whether an attribute value matches any value within a value list
EXISTS	Checks whether a subquery returns any rows
DISTINCT	Limits values to unique values
AGGREGATE FUNCTIONS	Used with SELECT to return mathematical summaries on columns
COUNT	Returns the number of rows with non-null values for a given column
MIN	Returns the minimum attribute value found in a given column
MAX	Returns the maximum attribute value found in a given column
SUM	Returns the sum of all values for a given column
AVG	Returns the average of all values for a given column

Data Definition Commands

- Examine simple database model and database tables that will form basis for many SQL examples
- Understand data environment



The Database Model



Creating the Database

- Following two tasks must be completed:
 - Create database structure
 - Create tables that will hold end-user data
- First task:
 - RDBMS creates physical files that will hold database
 - Tends to differ substantially from one RDBMS to another

The Database Schema

- Authentication
 - Process through which DBMS verifies that only registered users are able to access database
 - Log on to RDBMS using user ID and password created by database administrator
- Schema
 - Group of database objects—such as tables and indexes—that are related to each other

Data Types

- Data type selection is usually dictated by nature of data and by intended use
- Pay close attention to expected use of attributes for sorting and data retrieval purposes

Data Types (continued)

Some Common SQL Data Types

DATA TYPE	FORMAT	COMMENTS
Numeric	NUMBER(L,D)	The declaration NUMBER(7,2) indicates numbers that will be stored with two decimal places and may be up to six digits long, including the sign and the decimal place. Examples: 12.32, -134.99.
	INTEGER	May be abbreviated as INT. Integers are (whole) counting numbers, so they cannot be used if you want to store numbers that require decimal places.
	SMALLINT	Like INTEGER, but limited to integer values up to six digits. If your integer values are relatively small, use SMALLINT instead of INT.
	DECIMAL(L,D)	Like the NUMBER specification, but the storage length is a <i>minimum</i> specification. That is, greater lengths are acceptable, but smaller ones are not. DECIMAL(9,2), DECIMAL(9), and DECIMAL are all acceptable.
Character	CHAR(L)	Fixed-length character data for up to 255 characters. If you store strings that are not as long as the CHAR parameter value, the remaining spaces are left unused. Therefore, if you specify CHAR(25), strings such as "Smith" and "Katzenjammer" are each stored as 25 characters. However, a U.S. area code is always three digits long, so CHAR(3) would be appropriate if you wanted to store such codes.
	VARCHAR(L) or VARCHAR2(L)	Variable-length character data. The designation VARCHAR2(25) will let you store characters up to 25 characters long. However, VARCHAR will not leave unused spaces. Oracle users may use VARCHAR2 as well as VARCHAR.
Date	DATE	Stores dates in the Julian date format.

Creating Table Structures

- Use one line per column (attribute) definition
- Use spaces to line up attribute characteristics and constraints
- Table and attribute names are capitalized
- NOT NULL specification
- UNIQUE specification

Creating Table Structures (continued)

- Primary key attributes contain both a NOT NULL and a UNIQUE specification
- RDBMS will automatically enforce referential integrity for foreign keys
- Command sequence ends with semicolon

SQL Constraints

- NOT NULL constraint
 - Ensures that column does not accept nulls
- UNIQUE constraint
 - Ensures that all values in column are unique
- DEFAULT constraint
 - Assigns value to attribute when a new row is added to table
- CHECK constraint
 - Validates data when attribute value is entered

- When primary key is declared, DBMS automatically creates unique index
- Often need additional indexes
- Using CREATE INDEX command, SQL indexes can be created on basis of any selected attribute
- Composite index
 - Index based on two or more attributes
 - Often used to prevent data duplication

Data Manipulation Commands

- Adding table rows
- Saving table changes
- Listing table rows
- Updating table rows
- Restoring table contents
- Deleting table rows
- Inserting table rows with a select subquery



Adding Table Rows

- INSERT

- Used to enter data into table

- Syntax:

- INSERT INTO *columnname*
VALUES (*value1, value2, ... , valuen*);

Adding Table Rows (continued)

- When entering values, notice that:
 - Row contents are entered between parentheses
 - Character and date values are entered between apostrophes
 - Numerical entries are not enclosed in apostrophes
 - Attribute entries are separated by commas
 - A value is required for each column
- Use NULL for unknown values



Saving Table Changes

- Changes made to table contents are not physically saved on disk until, one of the following occurs:
 - Database is closed
 - Program is closed
 - COMMIT command is used
- Syntax:
 - COMMIT [WORK];
- Will permanently save any changes made to any table in the database

Listing Table Rows

- **SELECT**

- Used to list contents of table

- Syntax:

- `SELECT columnlist`
`FROM tablename;`

- *Columnlist* represents one or more attributes, separated by commas

- Asterisk can be used as wildcard character to list all attributes

Updating Table Rows

- UPDATE

- Modify data in a table

- Syntax:

- UPDATE *tablename*

- SET *columnname* = *expression* [, *columnname* = *expression*]

- [WHERE *conditionlist*];

- If more than one attribute is to be updated in row, separate corrections with commas

Restoring Table Contents

- **ROLLBACK**
 - Used to restore database to its previous condition
 - Only applicable if COMMIT command has not been used to permanently store changes in database
- **Syntax:**
 - ROLLBACK;
- **COMMIT and ROLLBACK only work with data manipulation commands that are used to add, modify, or delete table rows**

Deleting Table Rows

- DELETE

- Deletes a table row

- Syntax:

- DELETE FROM *tablename*
[WHERE *conditionlist*];

- WHERE condition is optional

- If WHERE condition is not specified, all rows from specified table will be deleted

Inserting Table Rows with a Select Subquery

- INSERT

- Inserts multiple rows from another table (source)
- Uses SELECT subquery
 - Query that is embedded (or nested) inside another query
 - Executed first
- Syntax:
 - INSERT INTO *tablename* SELECT *columnlist* FROM *tablename*;

Selecting Rows with Conditional Restrictions

- Select partial table contents by placing restrictions on rows to be included in output
 - Add conditional restrictions to SELECT statement, using WHERE clause
- Syntax:
 - SELECT *columnlist*
FROM *tablelist*
[WHERE *conditionlist*] ;

Selecting Rows with Conditional Restrictions (continued)

Example

Selecting Rows with Conditional Restrictions (continued)

SYMBOL	MEANING
=	Equal to
<	Less than
< =	Less than or equal to
>	Greater than
> =	Greater than or equal to
< > or !=	Not equal to

Arithmetic Operators: The Rule of Precedence

- Perform operations within parentheses
- Perform power operations
- Perform multiplications and divisions
- Perform additions and subtractions



Arithmetic Operators:

The Rule of Precedence (continued)

ARITHMETIC OPERATOR	DESCRIPTION
+	Add
-	Subtract
*	Multiply
/	Divide
^	Raise to the power of (Some applications use ** instead of ^.)

Special Operators

- **BETWEEN**

- Used to check whether attribute value is within a range

- **IS NULL**

- Used to check whether attribute value is null

- **LIKE**

- Used to check whether attribute value matches given string pattern

Special Operators (continued)

- **IN**
 - Used to check whether attribute value matches any value within a value list
- **EXISTS**
 - Used to check if subquery returns any rows

- All changes in table structure are made by using ALTER command
 - Followed by keyword that produces specific change
 - Following three options are available:
 - ADD
 - MODIFY
 - DROP



Changing a Column's Data Type

- ALTER can be used to change data type
- Some RDBMSs (such as Oracle) do not permit changes to data types unless column to be changed is empty



Changing a Column's Data Characteristics

- Use ALTER to change data characteristics
- If column to be changed already contains data, changes in column's characteristics are permitted if those changes do not alter the data type



Adding a Column

- Use ALTER to add column
 - Do not include the NOT NULL clause for new column

Dropping a Column

- Use ALTER to drop column
 - Some RDBMSs impose restrictions on the deletion of an attribute

Advanced Data Updates

Example

Copying Parts of Tables

- SQL permits copying contents of selected table columns so that the data need not be reentered manually into newly created table(s)
- First create the PART table structure
- Next add rows to new PART table using PRODUCT table rows



Adding Primary and Foreign Key Designations

- When table is copied, integrity rules do not copy, so primary and foreign keys need to be manually defined on new table
- User ALTER TABLE command
 - Syntax:
 - ALTER TABLE *tablename* ADD PRIMARY KEY(*fieldname*);
 - For foreign key, use FOREIGN KEY in place of PRIMARY KEY

Deleting a Table from the Database

- DROP
 - Deletes table from database
 - Syntax:
 - DROP TABLE *tablename*;



Advanced Select Queries

- SQL provides useful functions that can:
 - Count
 - Find minimum and maximum values
 - Calculate averages
- SQL allows user to limit queries to only those entries having no duplicates or entries whose duplicates may be grouped

Aggregate Functions

FUNCTION	OUTPUT
COUNT	The number of rows containing non-null values
MIN	The minimum attribute value encountered in a given column
MAX	The maximum attribute value encountered in a given column
SUM	The sum of all values for a given column
AVG	The arithmetic mean (average) for a specified column

Aggregate Functions (continued)

Example

Virtual Tables: Creating a View

- View is virtual table based on SELECT query
 - Can contain columns, computed columns, aliases, and aggregate functions from one or more tables
- Base tables are tables on which view is based
- Create view by using CREATE VIEW command



Virtual Tables: Creating a View (continued)

Example

Joining Database Tables

- Ability to combine (join) tables on common attributes is most important distinction between relational database and other databases
- Join is performed when data are retrieved from more than one table at a time
- Join is generally composed of an equality comparison between foreign key and primary key of related tables



Joining Tables with an Alias

- Alias can be used to identify source table
- Any legal table name can be used as alias
- Add alias after table name in FROM clause
 - FROM *tablename alias*

Summary

- SQL commands can be divided into two overall categories:
 - Data definition language commands
 - Data manipulation language commands
- The ANSI standard data types are supported by all RDBMS vendors in different ways
- Basic data definition commands allow you to create tables, indexes, and views



Summary (continued)

- DML commands allow you to add, modify, and delete rows from tables
- The basic DML commands are SELECT, INSERT, UPDATE, DELETE, COMMIT, and ROLLBACK
- INSERT command is used to add new rows to tables
- SELECT statement is main data retrieval command in SQL



Summary (continued)

- Many SQL constraints can be used with columns
- The column list represents one or more column names separated by commas
- WHERE clause can be used with SELECT, UPDATE, and DELETE statements to restrict rows affected by the DDL command



Summary (continued)

- Aggregate functions
 - Special functions that perform arithmetic computations over a set of rows
- ORDER BY clause
 - Used to sort output of SELECT statement
 - Can sort by one or more columns and use either an ascending or descending order
- Join output of multiple tables with SELECT statement

Summary (continued)

- Natural join uses join condition to match only rows with equal values in specified columns
- Right outer join and left outer join used to select rows that have no matching values in other related table

